deeplearning.ai

Deep Neural Networks

Deep L-layer Neural network

What is a deep neural network?

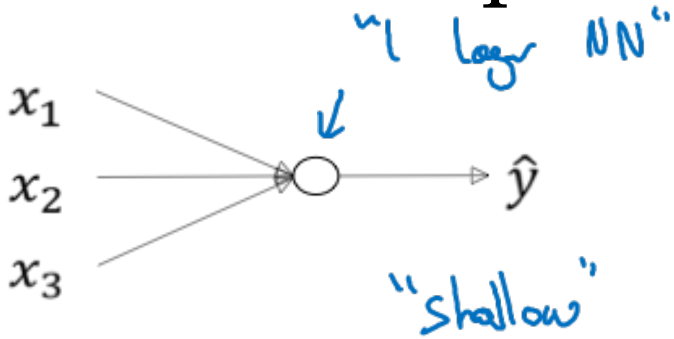

logistic regression

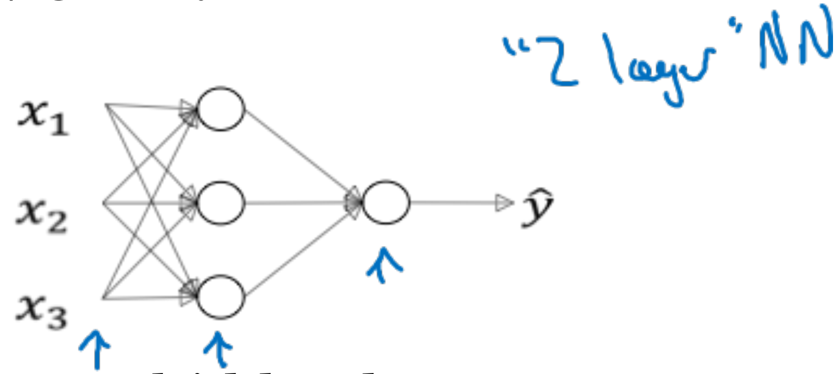

1 hidden layer

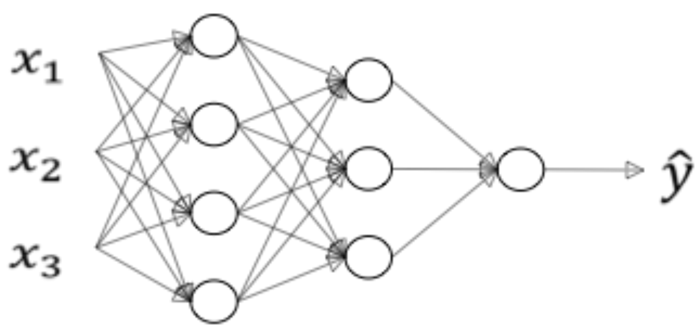

2 hidden layers

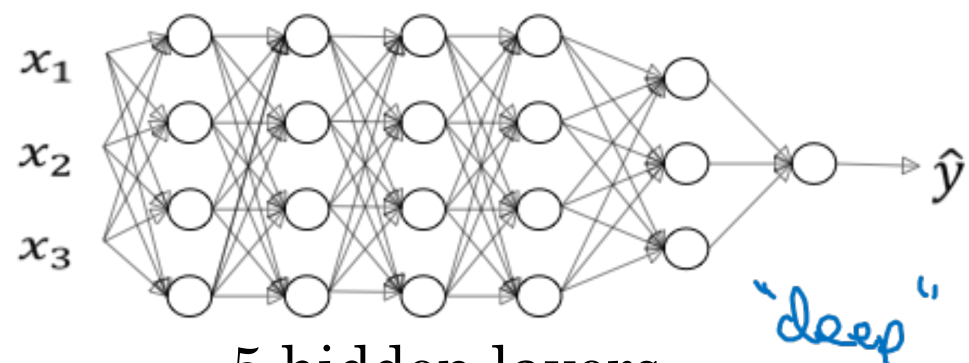

5 hidden layers

Andrew

Deep neural network notation

4 layer NN

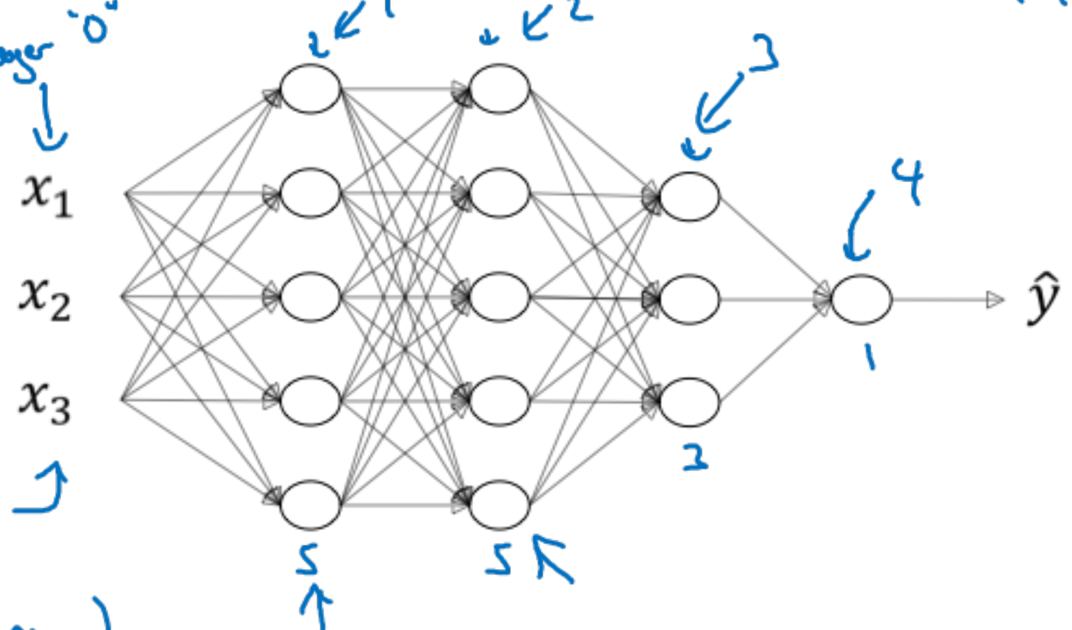

$X = a^{[0]}$

$\hat{y} = a^{[L]}$

$L = 4$ (#layers)

$n^{[l]}$ = #units in layer l

$n^{[1]} = 5,\ n^{[2]} = 5,\ n^{[3]} = 3,\ n^{[4]} = n^{[L]} = 1$

$n^{[0]} = n_x = 3$

$a^{[l]}$ = activations in layer l

$a^{[l]} = g^{[l]}(z^{[l]})$, $W^{[l]}$ = weights for $\underline{z^{[l]}}$

$b^{[l]}$

deeplearning.ai

Deep Neural Networks

Forward Propagation in a Deep Network

Forward propagation in a deep network

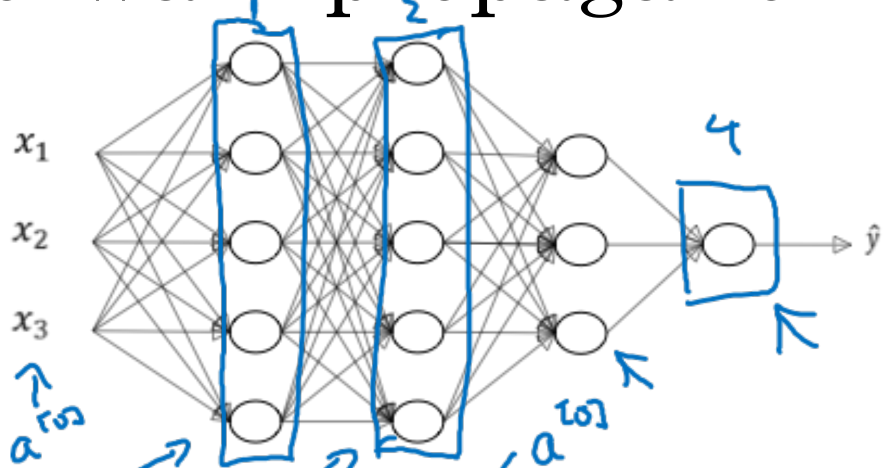

X: $z^{[1]} = W^{[1]} a^{[0]} + b^{[1]}$

$a^{[1]} = g^{[1]}(z^{[1]})$

$z^{[2]} = W^{[2]} a^{[1]} + b^{[2]}$

$a^{[2]} = g^{[2]}(z^{[2]})$

$\vdots$

$z^{[4]} = W^{[4]} a^{[3]} + b^{[4]}$, $a^{[4]} = g^{[4]}(z^{[4]}) = \hat{y}$

$\left[z^{[1](1)}, z^{[1](2)}, \dots z^{[1](m)} \right]$

$A^{[0]} = X$

$$\boxed{\begin{aligned} Z^{[l]} &= W^{[l]} A^{[l-1]} + b^{[l]} \\ A^{[l]} &= g^{[l]}(Z^{[l]}) \end{aligned}}$$

Vectorized:

$X = A^{[0]}$

$Z^{[1]} = W^{[1]} A^{[0]} + b^{[1]}$

$A^{[1]} = g^{[1]}(Z^{[1]})$

$Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]}$

$A^{[2]} = g^{[2]}(Z^{[2]})$

$\hat{Y} = g(Z^{[4]}) = A^{[4]}$

for $l = 1..4$

Deep Neural Networks

Getting your matrix dimensions right

deeplearning.ai

Parameters $W^{[l]}$ and $b^{[l]}$

$L = 5$

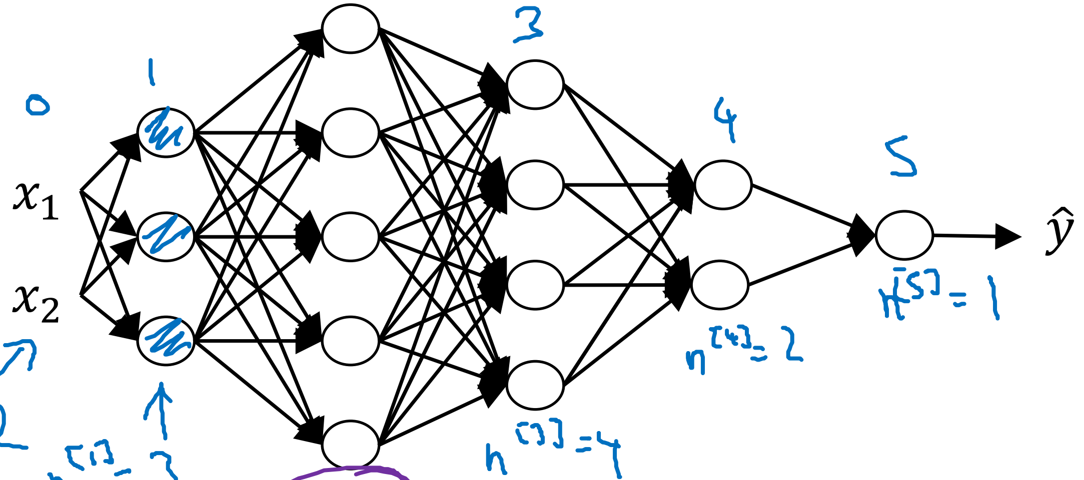

$n^{[0]} = n_x = 2$, $n^{[1]} = 3$, $n^{[2]} = 5$, $n^{[3]} = 4$, $n^{[4]} = 2$, $n^{[5]} = 1$

$z^{[l]} = g^{[l]}(a^{[l]})$

$a^{[l]}$

$W^{[l]} : (n^{[l]}, n^{[l-1]})$

$b^{[l]} : (n^{[l]}, 1)$

$dW^{[l]} : (n^{[l]}, n^{[l-1]})$

$db^{[l]} : (n^{[l]}, 1)$

$z^{[1]} = \boxed{W^{[1]} \cdot x} + \boxed{b^{[1]}}$

$(3,1) \leftarrow (3,2) \quad (2,1) \quad (3,1)$

$(n^{[1]}, 1) \quad (n^{[1]}, n^{[0]}) \quad (n^{[0]}, 1) \quad (n^{[1]}, 1)$

$W^{[1]} : (n^{[1]}, n^{[0]})$

$W^{[2]} : (5, 3) \quad (n^{[2]}, n^{[1]})$

$z^{[2]} = \boxed{W^{[2]} \cdot a^{[1]}} + \boxed{b^{[2]}}$

$(5,1) \quad (5,3) \quad (3,1) \quad (5,1) \quad (n^{[2]}, 1)$

$W^{[3]} : (4, 5)$

$W^{[4]} : (2, 4)$, $W^{[5]} : (1, 2)$

Andrew Ng

Vectorized implementation

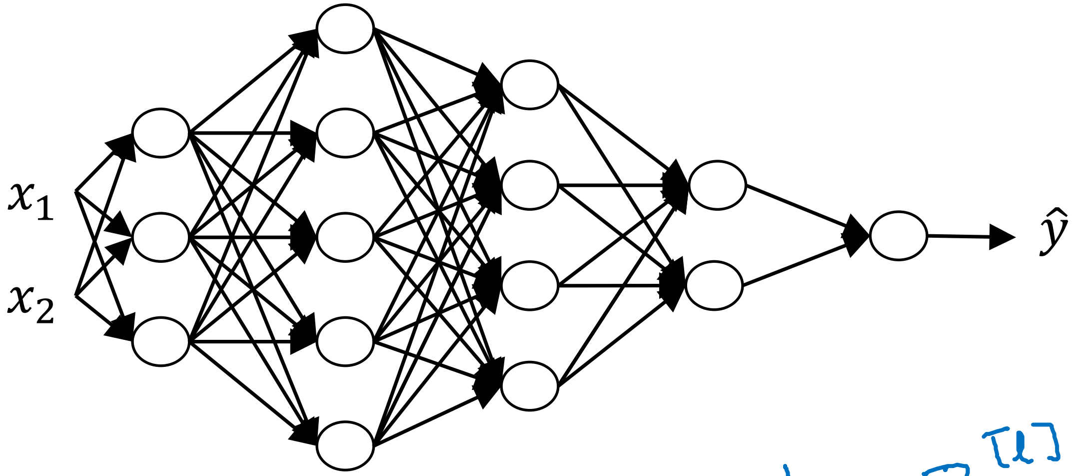

$z^{[1]} = W^{[1]} \cdot x + b^{[1]}$

$(n^{[1]}, 1) \quad (n^{[1]}, n^{[0]}) \quad (n^{[0]}, 1) \quad (n^{[1]}, 1)$

$\left[z^{[1](1)} \; z^{[1](2)} \; \cdots \; z^{[1](m)} \right]$

$Z^{[1]} = W^{[1]} \cdot X + b^{[1]}$

$(n^{[1]}, m) \quad (n^{[1]}, n^{[0]}) \quad (n^{[0]}, m) \quad (n^{[1]}, 1)$

$(n^{[1]}, m)$

$z^{[l]}, a^{[l]} : (n^{[l]}, 1)$

$Z^{[l]}, A^{[l]} : (n^{[l]}, m)$

$l = 0 \quad A^{[0]} = X = (n^{[0]}, m)$

$dZ^{[l]}, dA^{[l]} : (n^{[l]}, m)$

Andrew Ng

deeplearning.ai

Deep Neural Networks

Why deep representations?

Intuition about deep representation

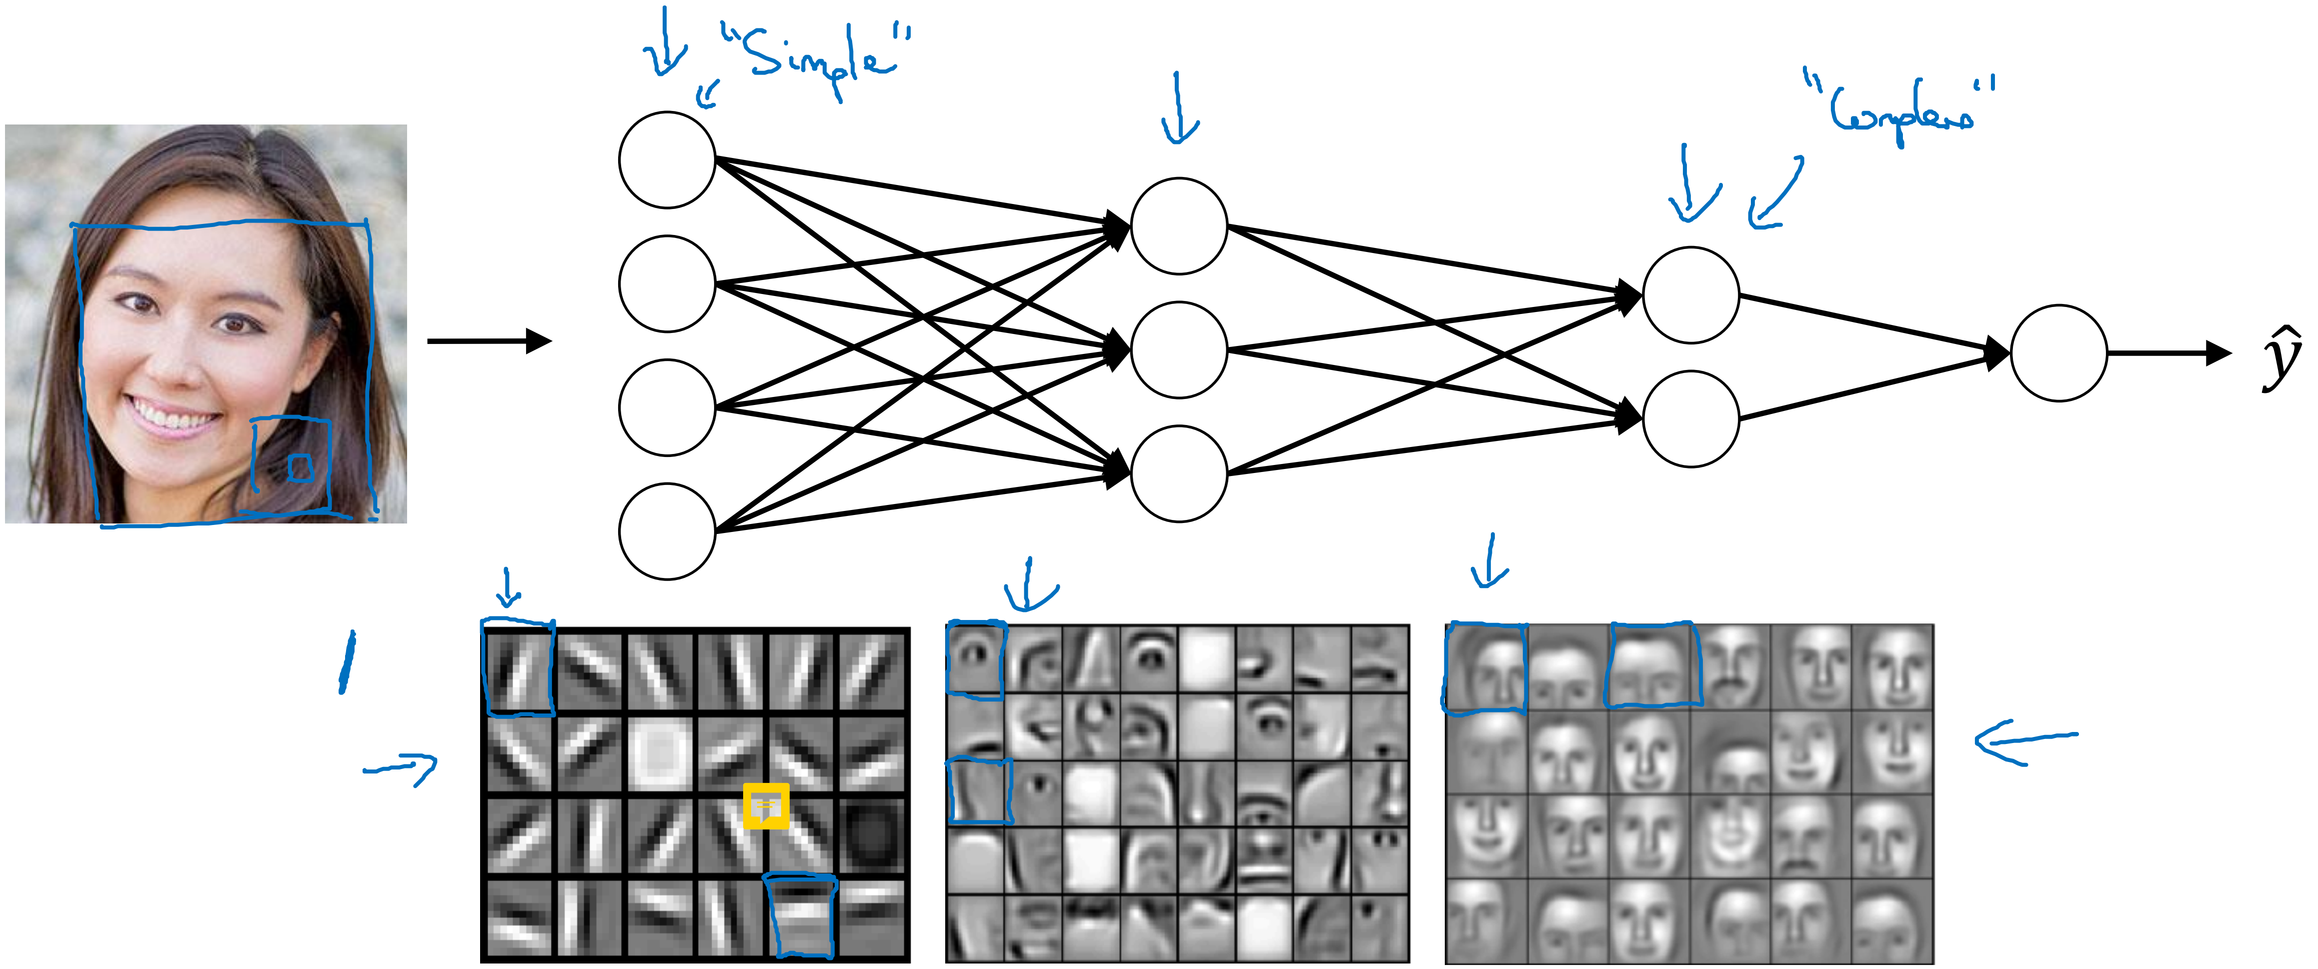

Audio → low level audio waveform → Phonemes C A T → Words → Sentence phrase

Circuit theory and deep learning

Informally: There are functions you can compute with a "small" L-layer deep neural network that shallower networks require exponentially more hidden units to compute.

$y = x_1 \text{ XOR } x_2 \text{ XOR } x_3 \text{ XOR} \ldots \text{ XOR } x_n$

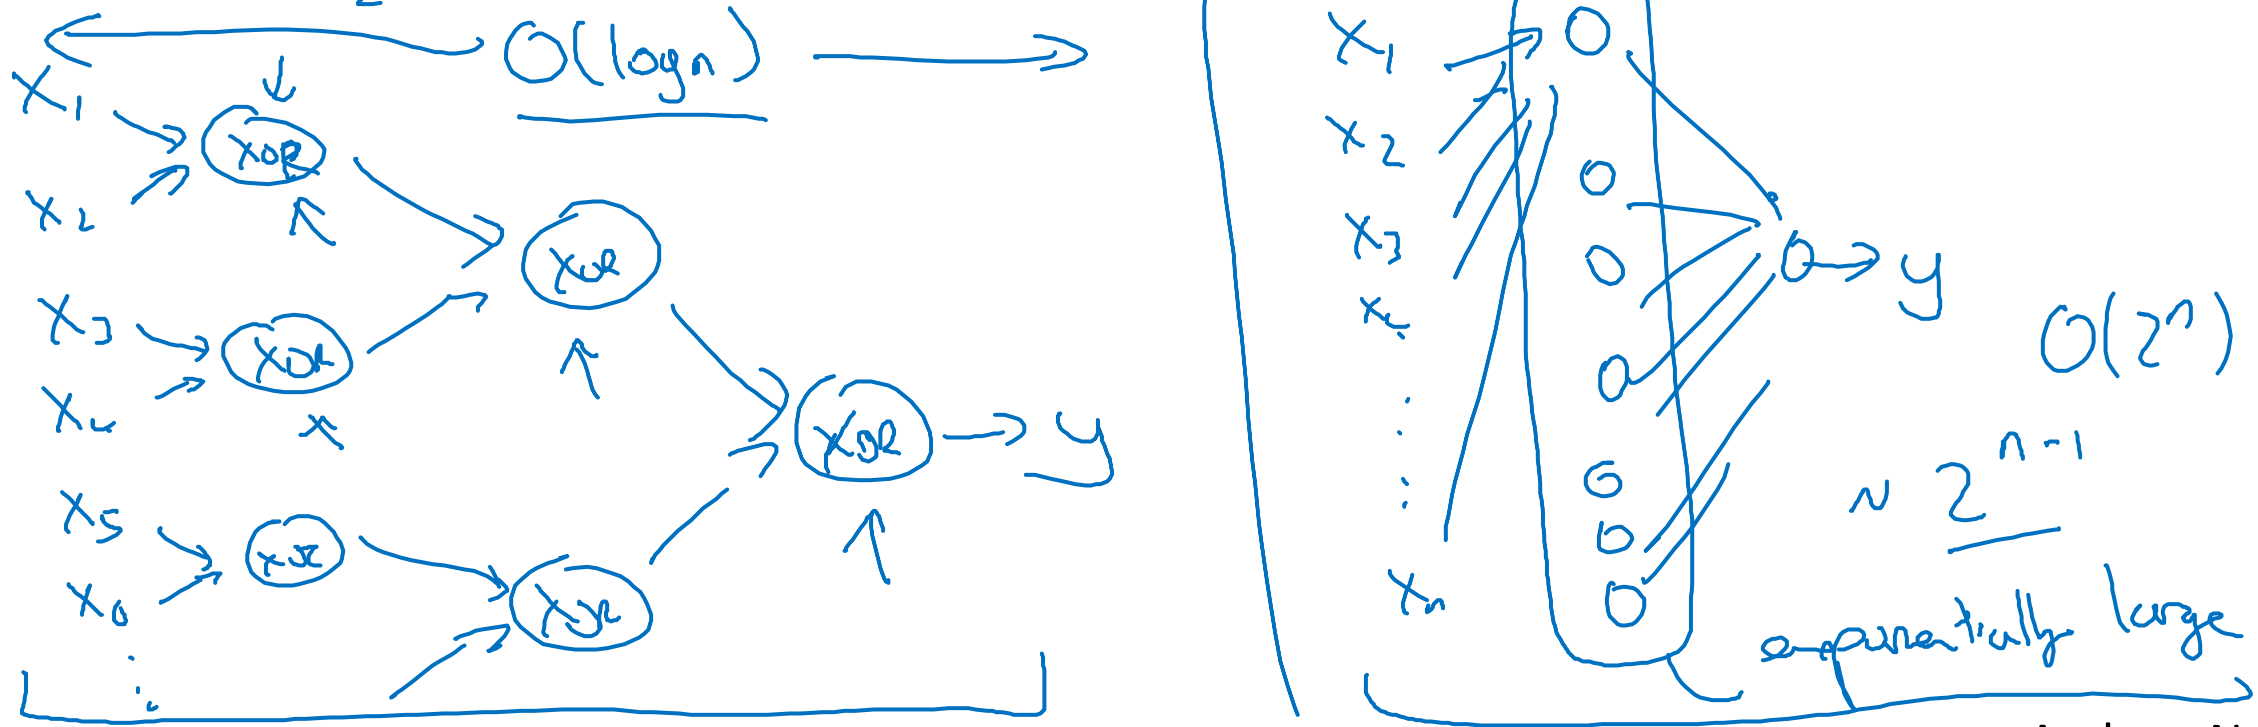

deeplearning.ai

Deep Neural Networks

Building blocks of deep neural networks

Forward and backward functions

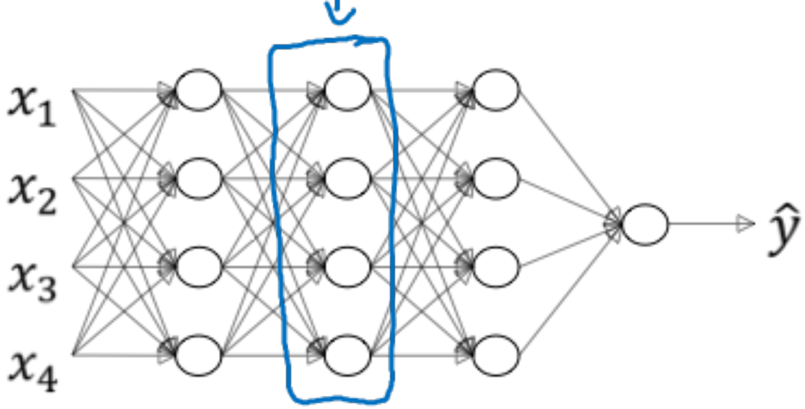

Layer l: $W^{[l]}, b^{[l]}$

→ Forward: Input $a^{[l-1]}$, output $a^{[l]}$

$\underline{z^{[l]}} = W^{[l]} \underline{a^{[l-1]}} + b^{[l]}$ cache $\underline{z^{[l]}}$

$\underline{a^{[l]}} = g^{[l]}(z^{[l]})$

→ Backward: Input $da^{[l]}$, cache $(z^{[l]})$ output $\underline{da^{[l-1]}}$, $dW^{[l]}$, $db^{[l]}$

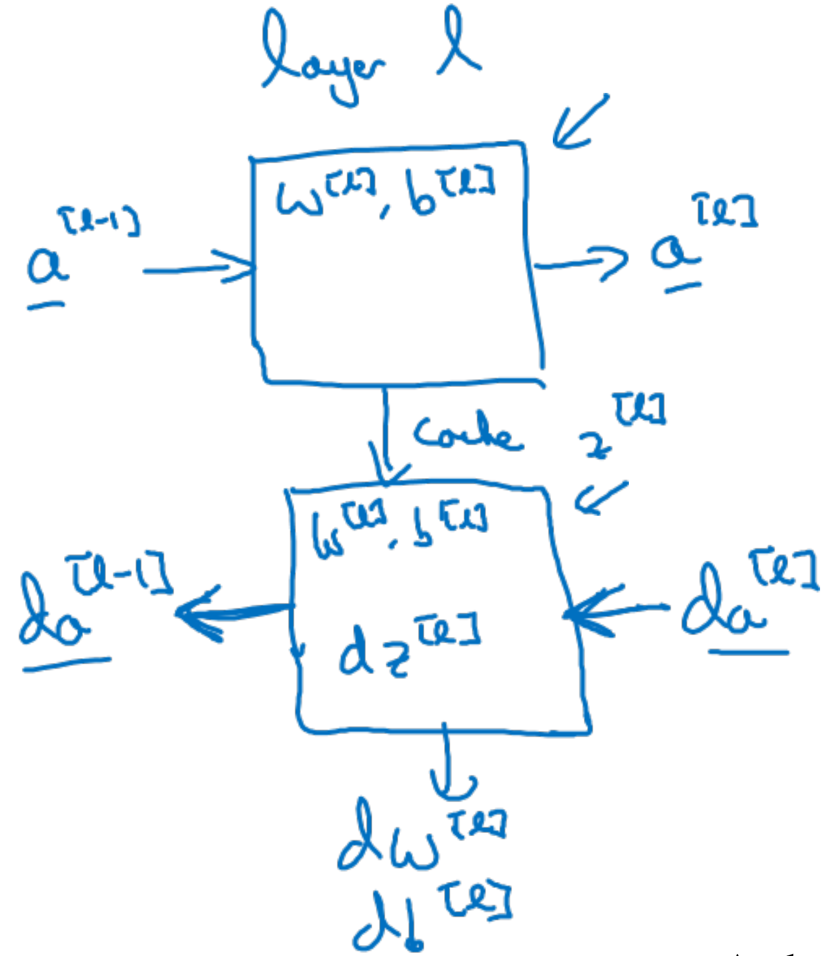

Forward and backward functions

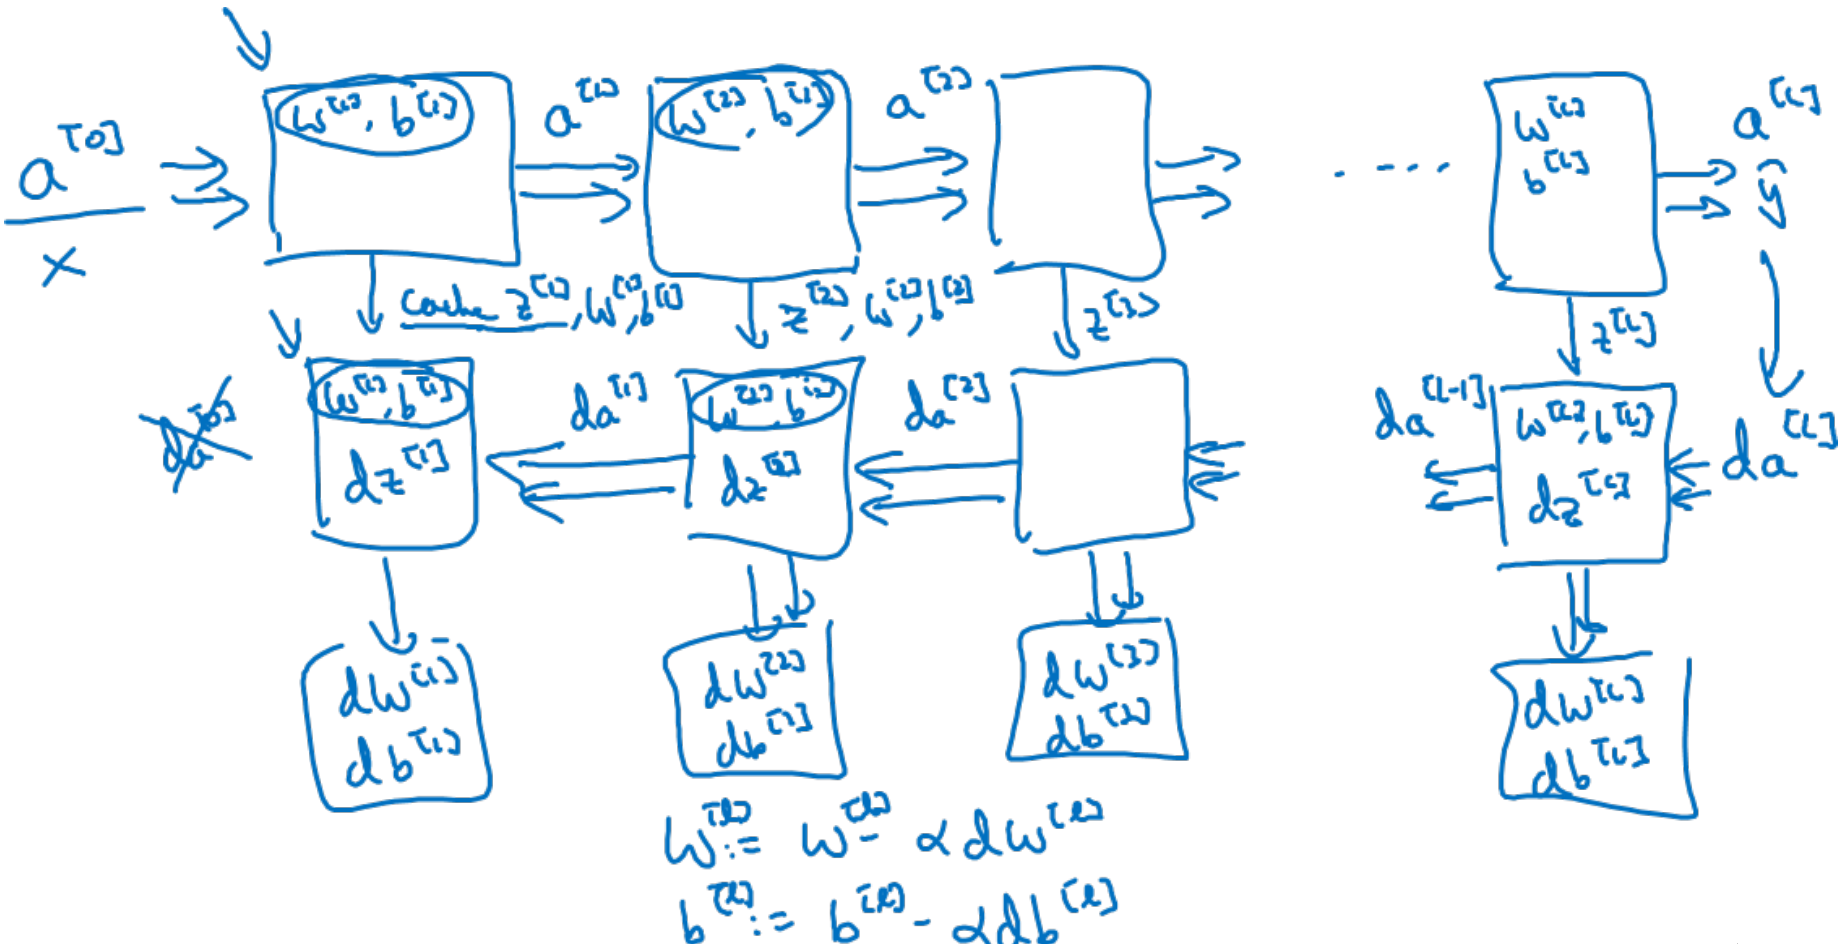

Andrew N

deeplearning.ai

Deep Neural Networks

Forward and backward propagation

Backward propagation for layer l

→ Input $da^{[l]}$

→ Output $da^{[l-1]}, dW^{[l]}, db^{[l]}$

$$dz^{[l]} = da^{[l]} * g^{[l]\prime}(z^{[l]})$$

$$dW^{[l]} = dz^{[l]} \cdot a^{[l-1]}$$

$$db^{[l]} = dz^{[l]}$$

$$da^{[l-1]} = W^{[l]T} \cdot dz^{[l]}$$

$$dz^{[l]} = W^{[l+1]T} dz^{[l+1]} * g^{[l]\prime}(z^{[l]})$$

$$dZ^{[l]} = dA^{[l]} * g^{[l]\prime}(Z^{[l]})$$

$$dW^{[l]} = \frac{1}{m} dZ^{[l]} \cdot A^{[l-1]T}$$

$$db^{[l]} = \frac{1}{m} np.sum(dZ^{[l]}, axis=1, keepdims=True)$$

$$dA^{[l-1]} = W^{[l]T} \cdot dZ^{[l]}$$

Summary

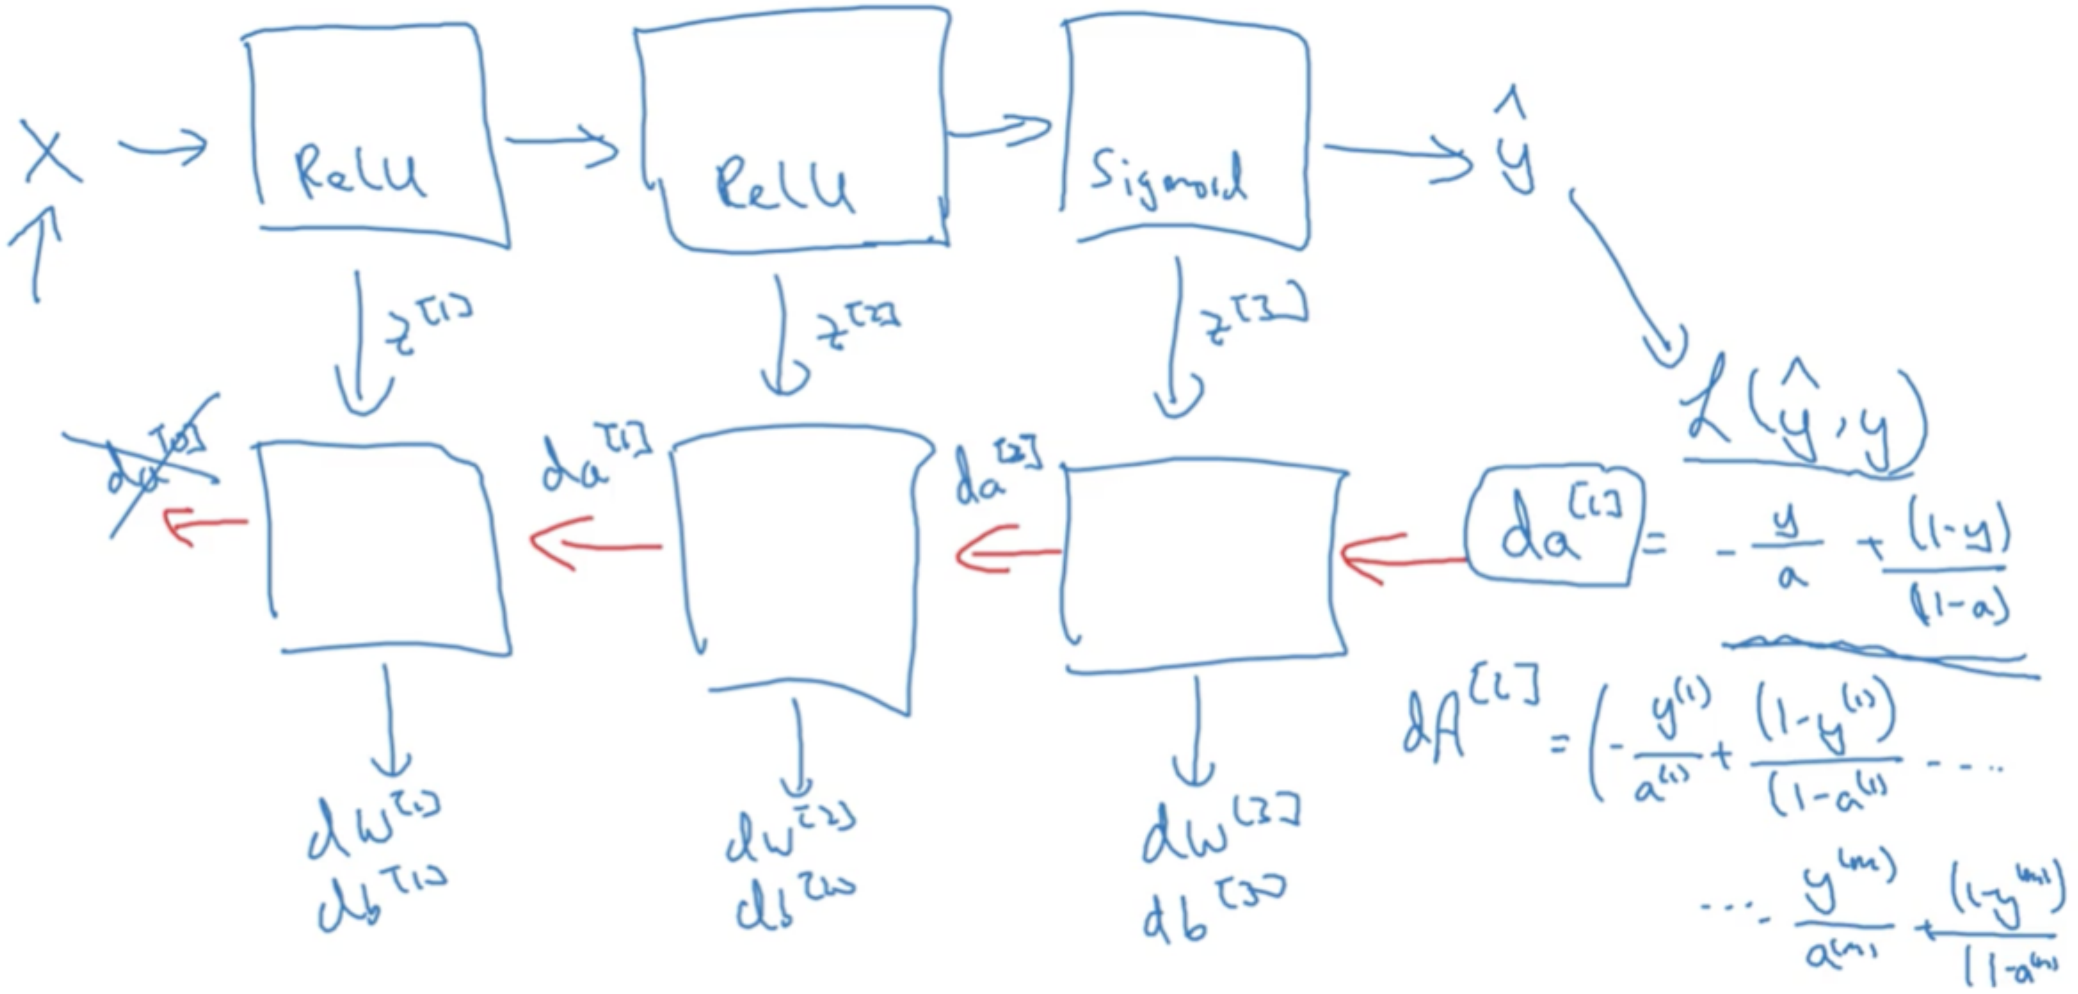

X
ReLU
ReLU
Sigmoid
ŷ
z[1]
z[2]
z[3]
L(ŷ, y)
da[0]
da[1]
da[2]
da[l] = -y/a + (1-y)/(1-a)
dA[l] = (-y(1)/a(1) + (1-y(1))/(1-a(1)) ... -y(m)/a(m) + (1-y(m))/(1-a(m)))
dW[1]
db[1]
dW[2]
db[2]
dW[3]
db[3]

deeplearning.ai

Deep Neural Networks

Parameters vs Hyperparameters

What are hyperparameters?

Parameters: $\underline{W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}, W^{[3]}, b^{[3]}}$ ...

Hyperparameters: learning rate α

#iterations

#hidden layers L

hidden units $n^{[1]}, n^{[2]}, \ldots$

choice of activation function

Later: Momentum, mini-batch size, regularizations, ...

Andrew Ng

Applied deep learning is a very empirical process

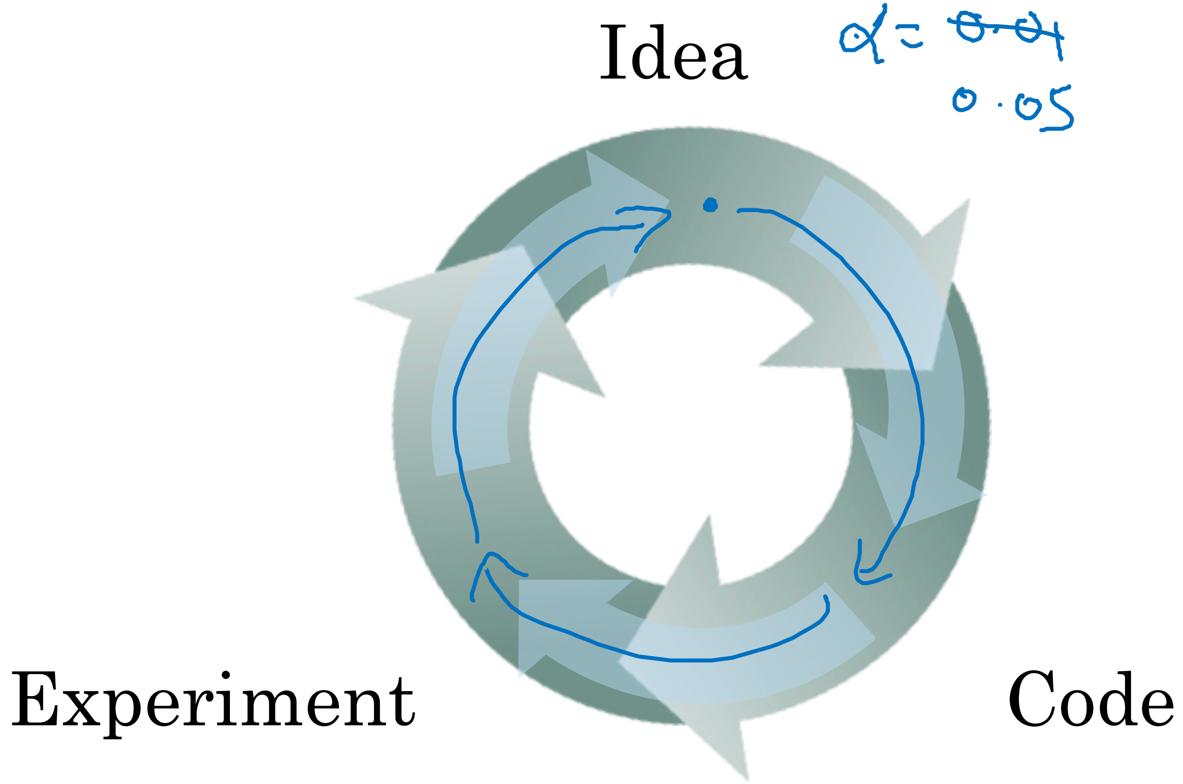

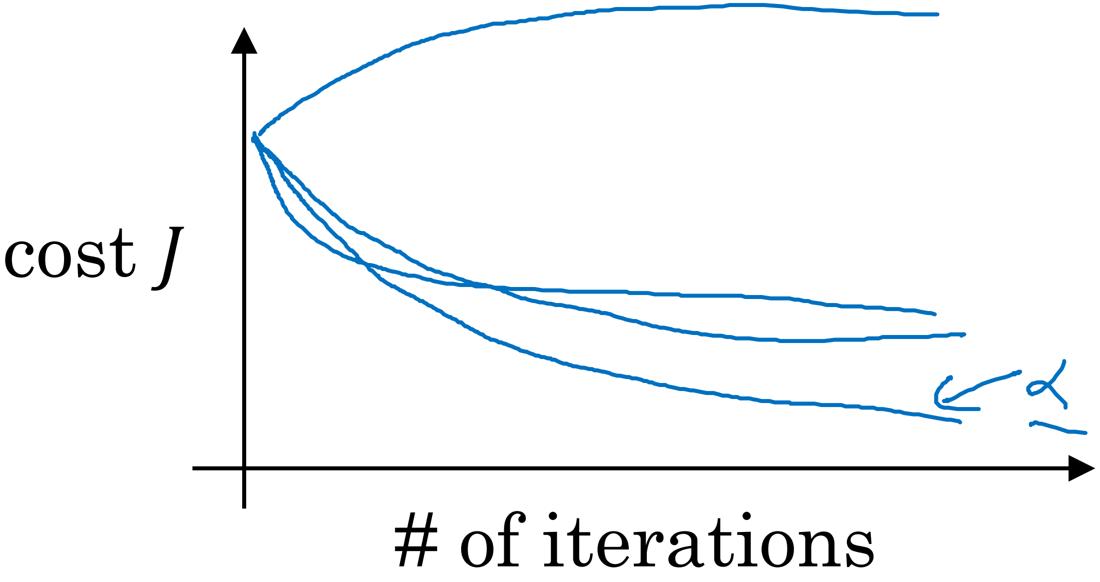

Vision, Speech, NLP, Ads, Search, Recommendation.

deeplearning.ai

Deep Neural Networks

What does this have to do with the brain?

Forward and backward propagation

$$Z^{[1]} = W^{[1]}X + b^{[1]}$$
$$A^{[1]} = g^{[1]}(Z^{[1]})$$
$$Z^{[2]} = W^{[2]}A^{[1]} + b^{[2]}$$
$$A^{[2]} = g^{[2]}(Z^{[2]})$$
$$\vdots$$
$$A^{[L]} = g^{[L]}(Z^{[L]}) = \hat{Y}$$

$$dZ^{[L]} = A^{[L]} - Y$$
$$dW^{[L]} = \frac{1}{m} dZ^{[L]} \cancel{A^{[L]^T}} \quad A^{[L-1]\,T}$$
$$db^{[L]} = \frac{1}{m} np.\mathrm{sum}(\mathrm{d}Z^{[L]}, axis = 1, keepdims = True)$$
$$dZ^{[L-1]} = \cancel{d}W^{[L]^T} dZ^{[L]} \cancel{g'^{[L]}}(Z^{[L-1]}) \quad * g'^{[L-1]}$$

Element wise Multiplication

$$\vdots$$
$$dZ^{[1]} = \cancel{d}W^{[2]} \; dZ^{[2]} * g'^{[1]}(Z^{[1]})$$
$$dW^{[1]} = \frac{1}{m} dZ^{[1]} A^{[0]^T}$$
$$db^{[1]} = \frac{1}{m} np.\mathrm{sum}(\mathrm{d}Z^{[1]}, axis = 1, keepdims = True)$$

"It's like the brain"

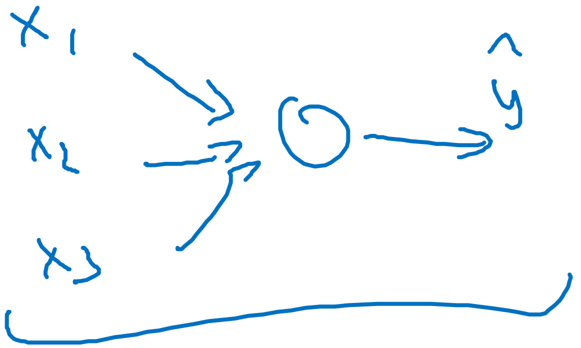

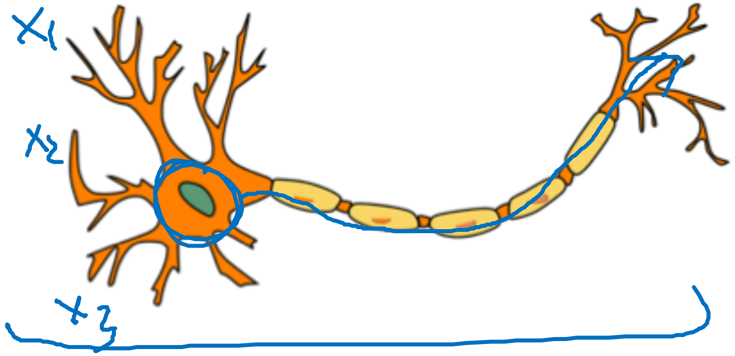

$X \longrightarrow y$

Andrew Ng